Improving language models by retrieving from trillions of tokens

Borgeaud et. al.

By Ilya Lasy

Outline

1. Motivation
2. Method
3. Related Work
4. Results
5. Limitations

Motivation

Decoupling "*Knowledge*" and "*Language information*" in Large Language Models.

- Efficiency (smaller amount of parameters with same performance)
- Interpretability
- Controllable generation
- Domain adaptation

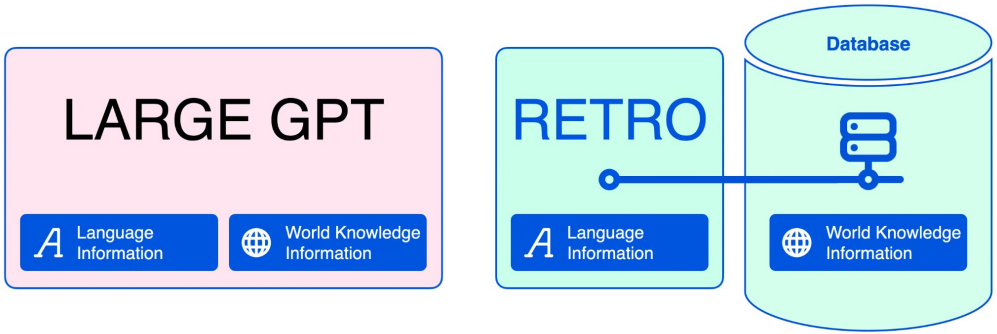

Method

RETRO (Retrieval-Enhanced Transformer)

1. Key-value database
2. Encoder-decoder
3. Chunked Cross-attention

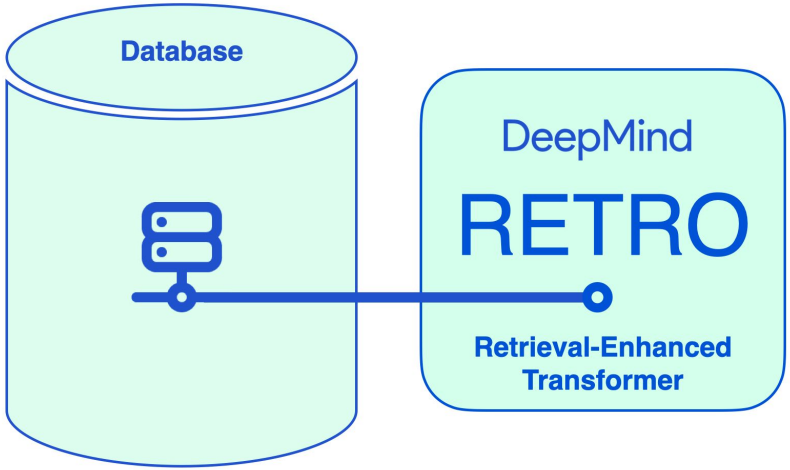

Dataset

MassiveText (Rae at. al.) - 5T tokens

Retrieval during *training*: 600B tokens

Retrieval during *evaluation*: up to 1.75T tokens

Source	Token count (M)	Documents (M)	Multilingual	Sampling frequency
Web	977,563	1,208	Yes	55%
Books	3,423,740	20	No	25%
News	236,918	398	No	10%
Wikipedia	13,288	23	Yes	5%
GitHub	374,952	143	No	5%

Retrieval Database

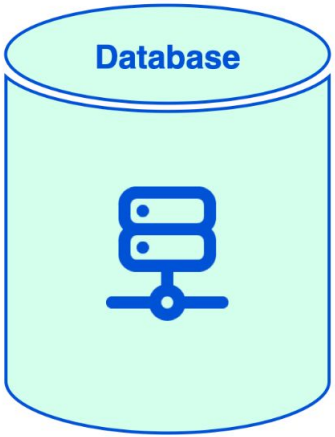

Key (BERT sentence embedding)	Value (text. neighbor and completion chunks. Each up to 64 tokens in length)	
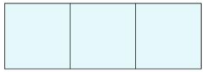	Dune is a 2021 American epic science fiction film directed by Denis Villeneuve	NEIGHBOR
	It is the first of a planned two-part adaptation of the 1965 novel by Frank Herbert	COMPLETION
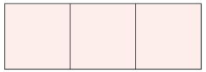	Dune is a 1965 science fiction novel by American author Frank Herbert	NEIGHBOR
	originally published as two separate serials in Analog magazine	COMPLETION
...	...	

Nearest neighbour retrieval

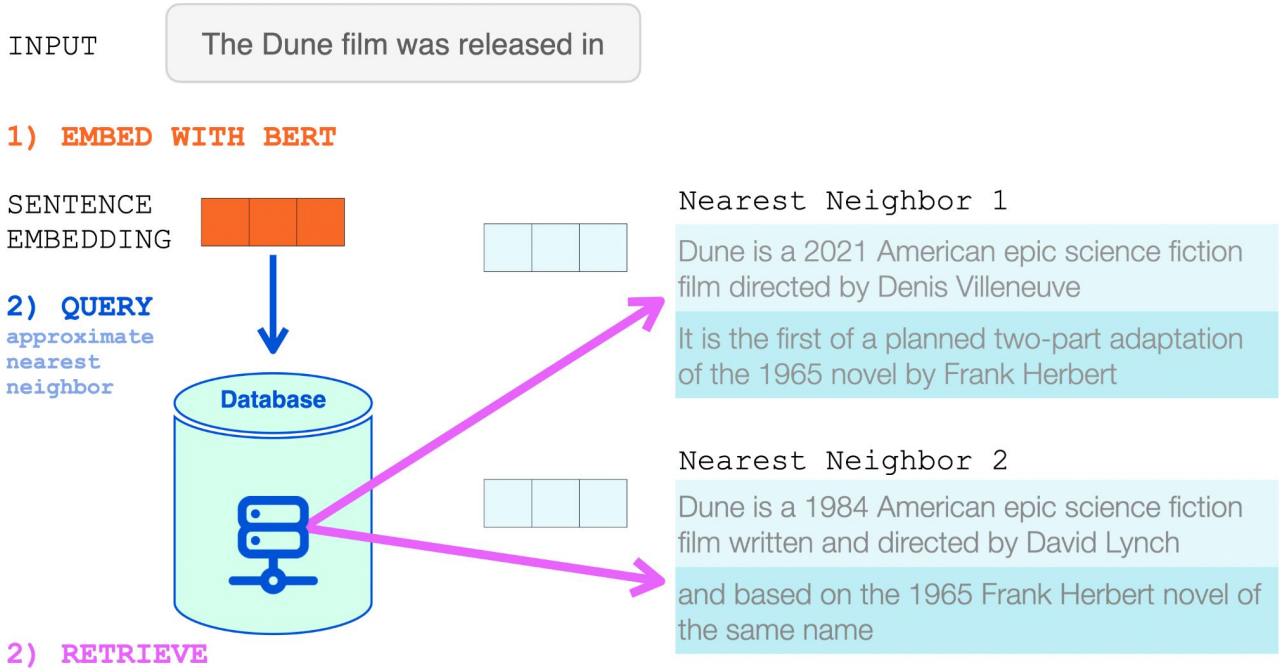

Faiss, SCaNN, etc.

High-level architecture

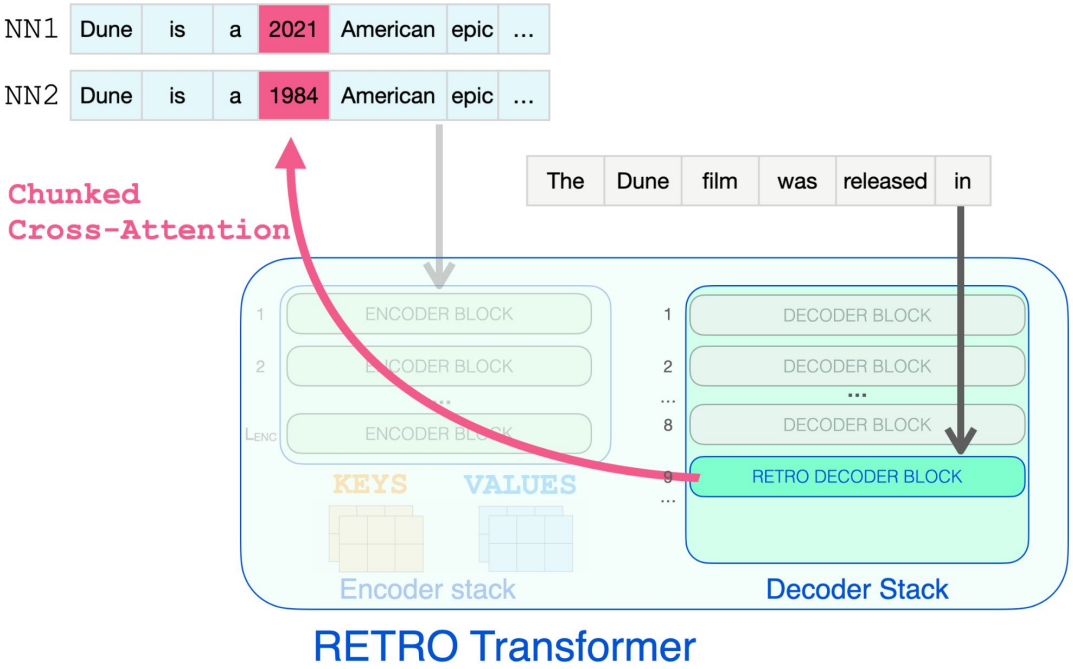

Encoder: non-autoregressive

Decoder: autoregressive

Detailed architecture

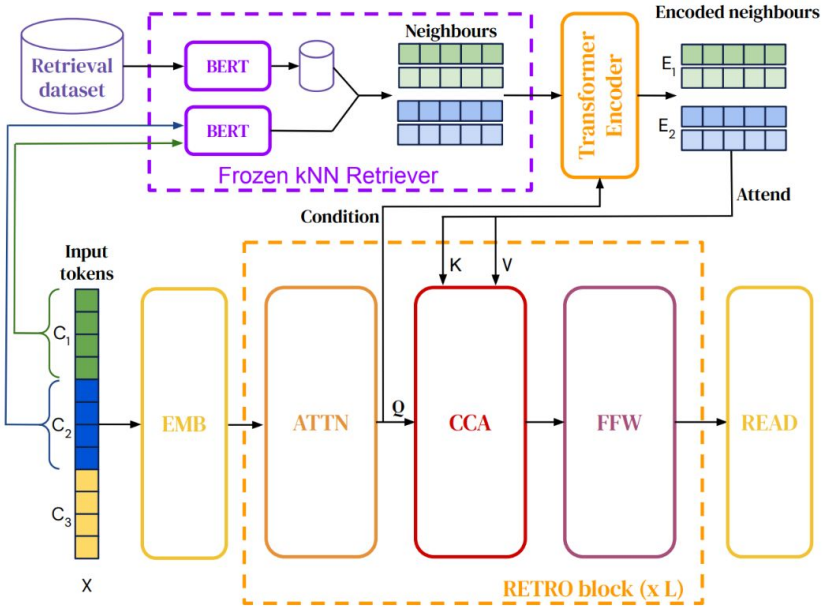

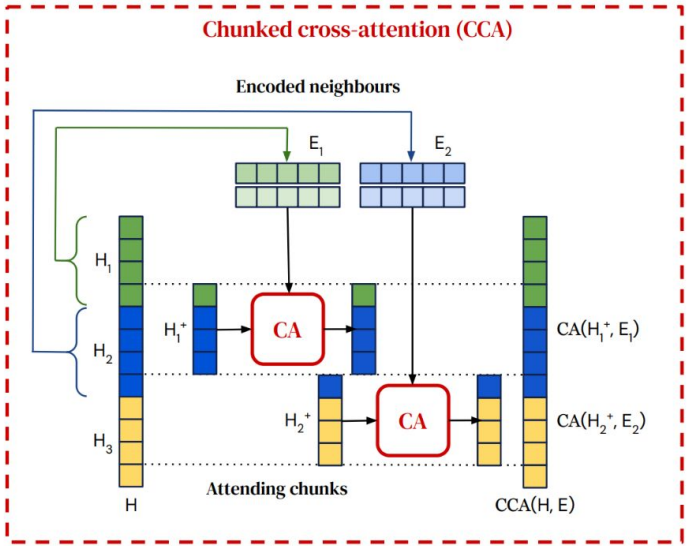

Decoder block

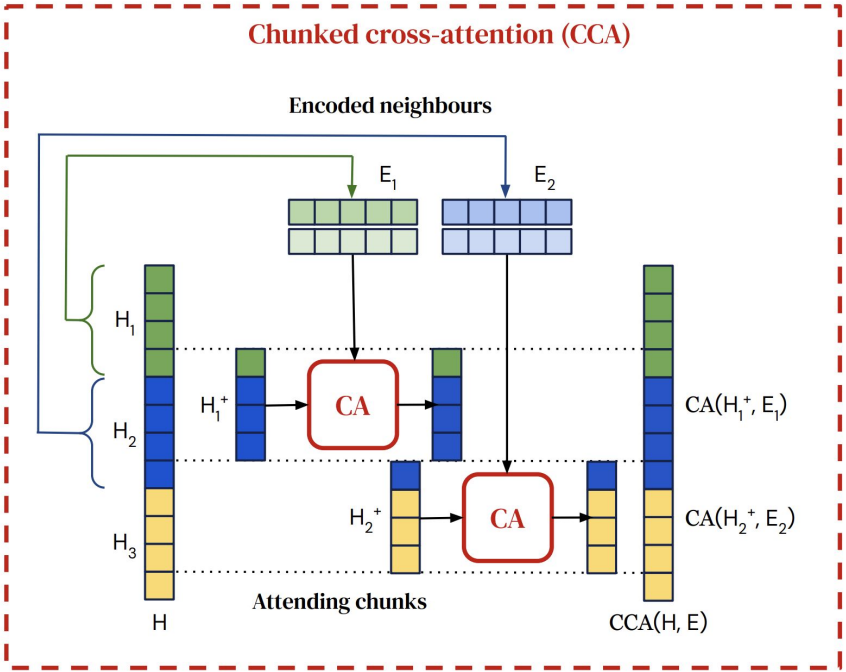

H_i - hidden activation of input tokens
E_i - encoded retrieved neighbors
CA *- cross-attention*

Retro-fitting

~10% of weights, ~3% of dataset

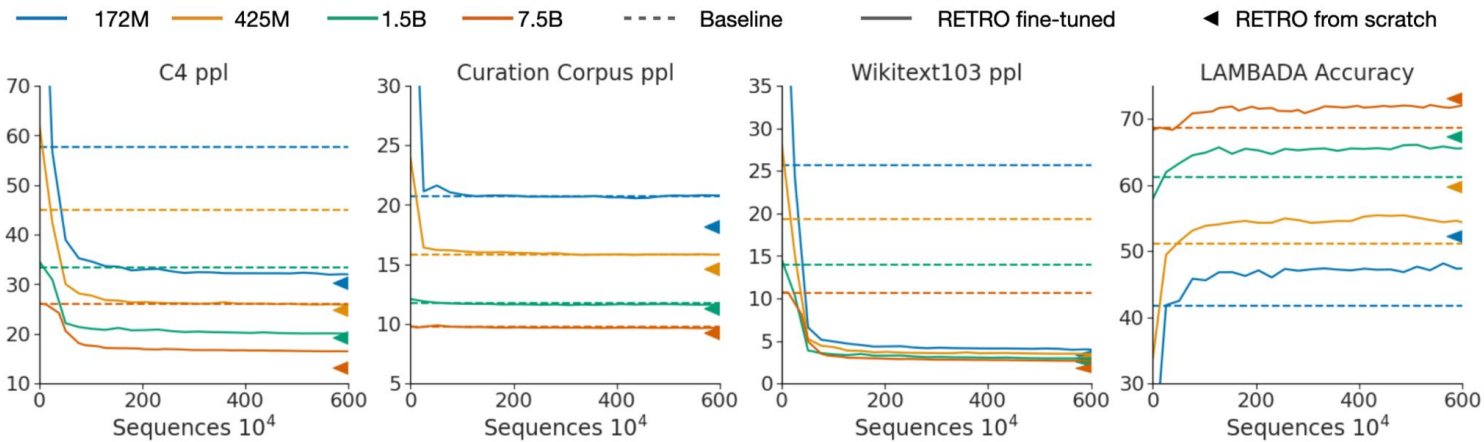

Figure: Retro-fitting a baseline transformer

Related work

	# Retrieval tokens	Granularity	Retriever training	Retrieval integration
Continuous Cache	$O(10^3)$	Token	Frozen (LSTM)	Add to probs
kNN-LM	$O(10^9)$	Token	Frozen (Transformer)	Add to probs
SPALM	$O(10^9)$	Token	Frozen (Transformer)	Gated logits
DPR	$O(10^9)$	Prompt	Contrastive proxy	Extractive QA
REALM	$O(10^9)$	Prompt	End-to-End	Prepend to prompt
RAG	$O(10^9)$	Prompt	Fine-tuned DPR	Cross-attention
FID	$O(10^9)$	Prompt	Frozen DPR	Cross-attention
EMDR^2	$O(10^9)$	Prompt	End-to-End (EM)	Cross-attention
RETRO (ours)	$O(10^{12})$	**Chunk**	**Frozen (BERT)**	**Chunked cross-attention**

Table: Comparison of Retro with existing retrieval approaches.

Evaluation metric: bits-per-bytes

Bits-per-byte (bpb) measures the average number of bits required to predict the next token in a sequence, which is usually a byte.

The **lower** the bpb value - the more **efficient** the model in making predictions, i.e., it requires fewer bits to predict the next token.

$$L \times \log_2(e)$$

where ***L*** is the cross-entropy loss

Results

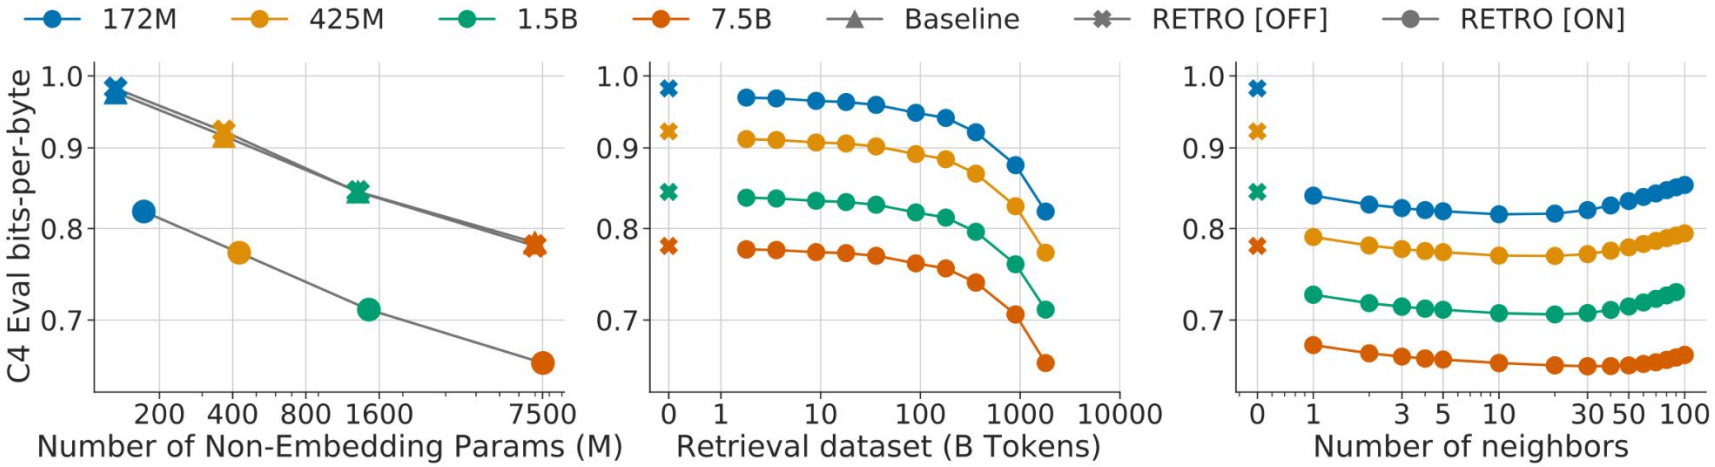

Figure: Scaling of Retro

Results

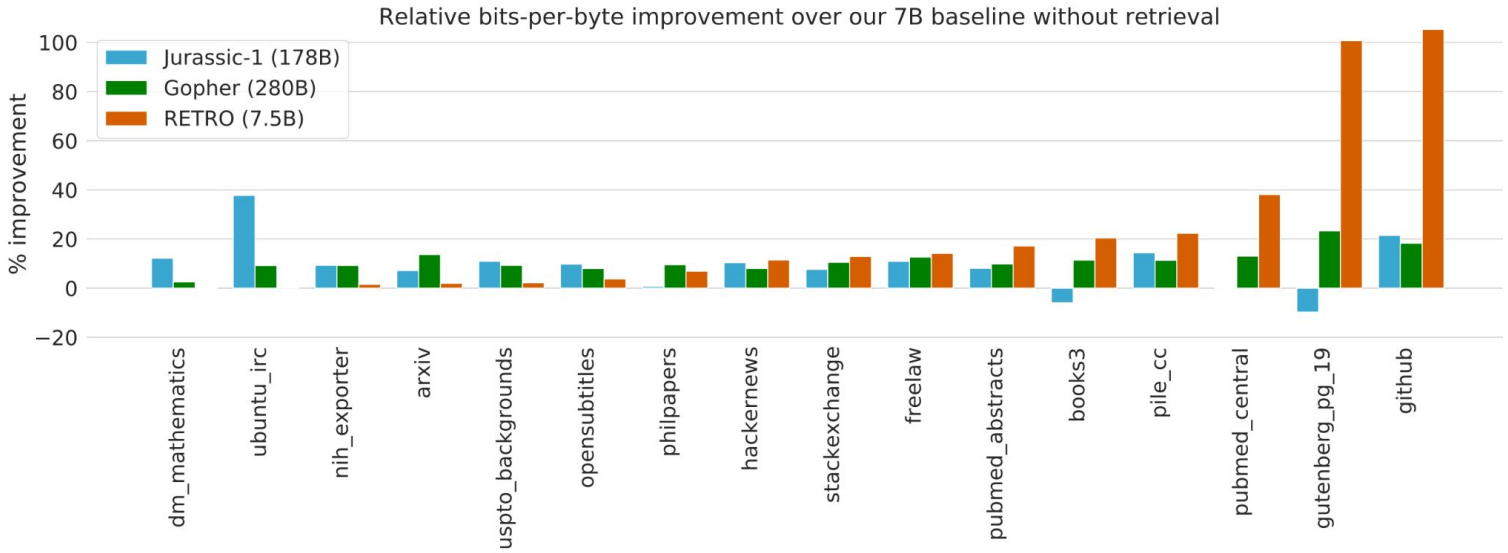

Figure: comparison to non-retrieval LLM

Results

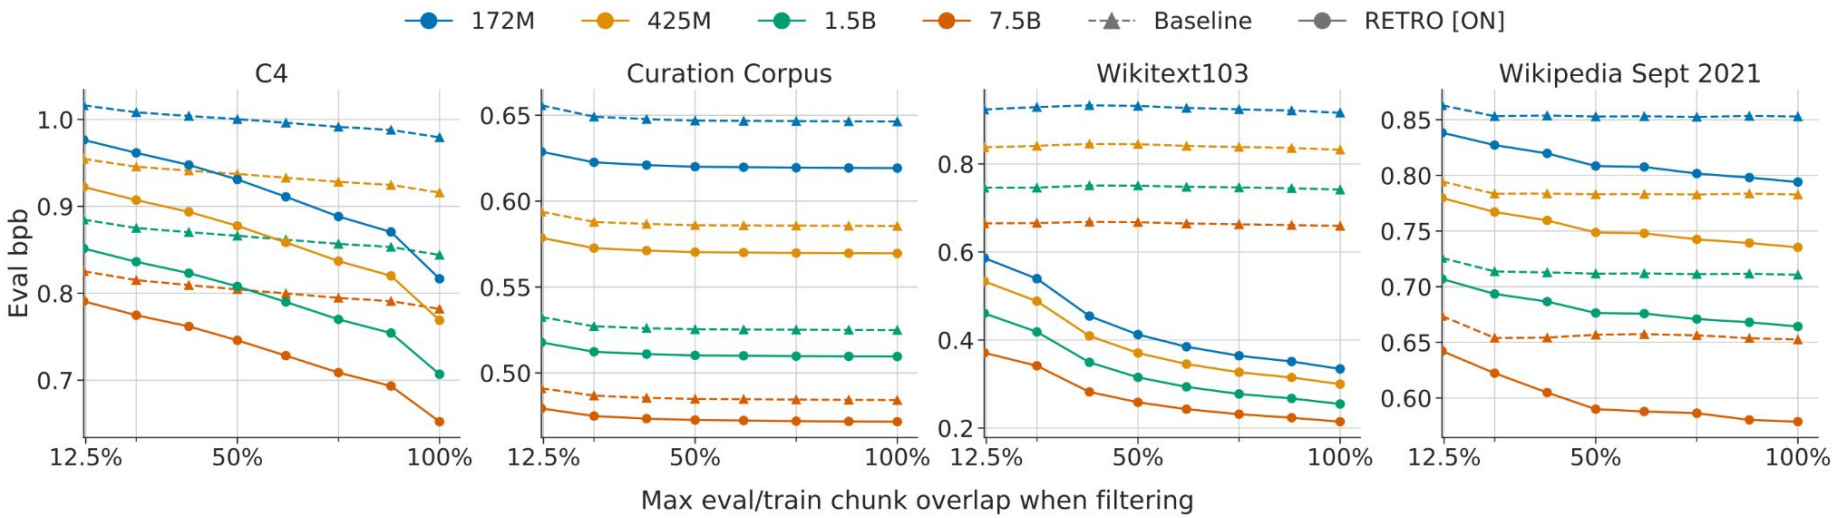

Figure: Relating retrieval performance to dataset leakage

Results - sample

Prompt and sample of RETRO[OFF]	Prompt and sample of RETRO[ON] colored by LCP with RET(C_u−1) LCP = 0, 1, 2, 3,4,⩾ 5	$[N_u^1, F_u^1]$ colored by LCP with C_{u+1} LCP = 0, 1, 2, 3,4,⩾ 5	$[N_u^2, F_u^2]$ colored by LCP with C_{u+1} LCP = 0, 1, 2, 3,4,⩾ 5
Beavers are interesting animals that live near rivers. They build	Beavers are interesting animals that live near rivers. They build	.Beavers build their lodges in pon ds they have created in wooded areas .Like many things in nature, there is a connection between creatures i n the wild.Beaver ponds cause tree s to drown, but the dead trees attra ct the great blue heron, which often return year after year. Over time, a beaver pond can attract more than 50 nests in a colony, called a rooke ry.An example of this can be found in the large pond off Bradford Road at Carter Fields near the Boxford l ine.Chris Leahy, an expert with th e Massachusetts Audubon Society who wrote	naw them into smaller sections and d rag them into the water.Engineers of the Pond Beavers are interesting animals because they change the hab itat in which they live. Beavers do this by blocking up streams to creat e ponds. Then they build their homes , called lodges, in these ponds. Bea vers' bodies make them well-suited f or underwater building Special muscl es close off their noses, ears, and throats to keep the water out. Beave rs' broad tails act like rudders for steering. Their two very large, ora nge front teeth are used to gnaw dow n trees. They begin building their d am
dams to create ponds. Frogs are am phibians, so they can live in both l and and water. They have great camou flage to hide from predators. The G olden Retriever, or Golden, is a ver y big dog. They are very strong, and have a thick coat to help them live in	their houses called beaver dams in the riverbeds. They also live on lan d.Beavers use their strong teeth an d strong jaws to cut down trees and branches to build their homes. They also use their teeth and jaws to che w up food. Beavers use their big, fl at tails to swim. They use	, then they mean that you are very b usy. Beavers swim easily in streams, picking up rocks and sticks to buil d their dams. They gnaw at trees wit h their big front teeth to cut them down. Then they use parts of the tre es to build their houses.Beavers ar e clever builders. They know exactly what they need to build their beave r dams. They use mud from the stream to make their dams stay together. T hey use their tails to pat down the mud.Beavers put a snug room at the top of their dams for their babies. They store their food underwater. Be avers eat the bark from the	ar-like tail, and two protruding tee th that are strong enough to gnaw do wn trees. The beaver uses trees, bra nches, and mud to build dams across rivers and streams. These dams creat e a deep pool of water in which the beaver builds its home. A beaver hom e is called a lodge. A baby beaver o r "kit" remains in the family lodge until the age of two.Beaver fur, kn own as pelt, was once highly popular as a trim for hats and coats. How m ight the popularity of beaver fur co ntributed to the colonization of New York?wwwWw

Limitations

- Still Transformer - 100% decoupling is not achieved.
- Retrieved content length is still limited.
- Knowledge base is just unstructured text.
- What is “world knowledge”?
- What is “language information”?

References

- *Borgeaud et.al.* 2022. Improving language models by retrieving from trillions of tokens
- *Wang et.al.* 2023. Shall We Pretrain Autoregressive Language Models with Retrieval? A Comprehensive Study
- *Wang et.al.* 2023. InstructRetro: Instruction Tuning post Retrieval-Augmented Pretraining
- The Illustrated Retrieval Transformer